

CSCI547 Machine Learning

Homework 1

Zachary Falkner
Department of Computer Science
University of Montana

March 6, 2018

1 Linear Regression

1A

$$(X^T X + \gamma I_N) \vec{W} = X^T Y$$

$$\vec{w} = [w_0, w_1, w_2, w_3]$$

$$X =$$

$$\begin{bmatrix} 1 & x_0 & x_0^2 & x_0^3 \\ 1 & x_1 & x_1^2 & x_1^3 \\ \dots & \dots & \dots & \dots \\ 1 & x_m & x_m^2 & x_m^3 \end{bmatrix}$$

$$\mathcal{J} = \frac{1}{2} (\sum_{i=1}^m (y_i - w_0 - w_1 x_i - w_2 x_i^2 - w_3 x_i^3))^2$$

$$\frac{\partial \mathcal{J}}{\partial w_0} = - \sum_{i=1}^m (y_i - w_0 - w_1 x_i - w_2 x_i^2 - w_3 x_i^3)$$

$$\frac{\partial \mathcal{J}}{\partial w_1} = - \sum_{i=1}^m (x_i (y_i - w_0 - w_1 x_i - w_2 x_i^2 - w_3 x_i^3)) + 2\gamma w_1$$

$$\frac{\partial \mathcal{J}}{\partial w_2} = - \sum_{i=1}^m (x_i^2 (y_i - w_0 - w_1 x_i - w_2 x_i^2 - w_3 x_i^3)) + 2\gamma w_2$$

$$\frac{\partial \mathcal{J}}{\partial w_3} = - \sum_{i=1}^m (x_i^3 (y_i - w_0 - w_1 x_i - w_2 x_i^2 - w_3 x_i^3)) + 2\gamma w_3$$

$$\frac{\partial \mathcal{J}}{\partial w_0} = \sum_{i=1}^m (-y_i + w_0 + w_1 x_i + w_2 x_i^2 + w_3 x_i^3)$$

$$\frac{\partial \mathcal{J}}{\partial w_1} = \sum_{i=1}^m (-y_i x_i + w_0 x_i - w_1 x_i^2 - w_2 x_i^3 - w_3 x_i^4) + 2\gamma w_1$$

$$\frac{\partial \mathcal{J}}{\partial w_2} = \sum_{i=1}^m (-y_i x_i^2 + w_0 x_i^2 + w_1 x_i^3 + w_2 x_i^4 + w_3 x_i^5) + 2\gamma w_2$$

$$\frac{\partial \mathcal{J}}{\partial w_3} = \sum_{i=1}^m (-y_i x_i^3 + w_0 x_i^3 + w_1 x_i^4 + w_2 x_i^5 + w_3 x_i^6) + 2\gamma w_3$$

$$\frac{\partial \mathcal{J}}{\partial w_0} = 0 = \sum_{i=1}^m (-y_i + w_0 + w_1 x_i + w_2 x_i^2 + w_3 x_i^3)$$

$$\frac{\partial \mathcal{J}}{\partial w_1} = 0 = \sum_{i=1}^m (-y_i x_i + w_0 x_i - w_1 x_i^2 - w_2 x_i^3 - w_3 x_i^4) + 2\gamma w_1$$

$$\frac{\partial \mathcal{J}}{\partial w_2} = 0 = \sum_{i=1}^m (-y_i x_i^2 + w_0 x_i^2 + w_1 x_i^3 + w_2 x_i^4 + w_3 x_i^5) + 2\gamma w_2$$

$$\frac{\partial \mathcal{J}}{\partial w_3} = 0 = \sum_{i=1}^m (-y_i x_i^3 + w_0 x_i^3 + w_1 x_i^4 + w_2 x_i^5 + w_3 x_i^6) + 2\gamma w_3$$

$$\sum_{i=1}^m y_i = \sum_{i=1}^m (+w_0 + w_1 x_i + w_2 x_i^2 + w_3 x_i^3)$$

$$\sum_{i=1}^m y_i = \frac{\sum_{i=1}^m (w_0 x_i - w_1 x_i^2 - w_2 x_i^3 - w_3 x_i^4) + 2\gamma w_1}{x_i}$$

$$\sum_{i=1}^m y_i = \frac{\sum_{i=1}^m (w_0 x_i^2 + w_1 x_i^3 + w_2 x_i^4 + w_3 x_i^5) + 2\gamma w_2}{x_i^2}$$

$$\sum_{i=1}^m y_i = \frac{\sum_{i=1}^m (w_0 x_i^3 + w_1 x_i^4 + w_2 x_i^5 + w_3 x_i^6) + 2\gamma w_3}{x_i^3}$$

1B

```
import pandas as pd
import numpy as np
import matplotlib as mpl
import matplotlib.pyplot as plt

#         degree - an arbitrary degree of polynomial
#         gamma - regularization strength (constant)
#         dataset - a pandas dataframe
#
def fit_with_l2(degree, gamma, dataset, test_set):
    #data and target from pandas frame
    x = dataset[0].as_matrix().astype(float)
    y = dataset[1].as_matrix().astype(float)

    x_test = test_set[0].as_matrix().astype(float)
    y_test = test_set[1].as_matrix().astype(float)

    xhat = np.linspace(x.min(), x.max(), 200)
    Xhat = np.vander(xhat, N=degree + 1, increasing=True)
    x = 2*(dataset[0] - x.min())/(x.max()-x.min()) - 1 #this is x_tilde
    x_test = 2*(test_set[0] - x.min())/(x.max()-x.min()) - 1

    # Vandermonde matrix (design matrix)
    X = np.vander(x, degree+1, increasing=True)

    # Identity matrix for regularization
    Eye = np.eye(X.shape[1])
    #Eye hat, to not regularize bias
    Eye[0, 0] = 0

    # Solve for weight's set (params) (training)
    w = np.linalg.solve(np.dot(X.T, X) + gamma*Eye, np.dot(X.T, y))
```

```

yhat = np.dot(Xhat,w) # 250x2, 16x1

X_test = np.vander(x_test, N=degree+1, increasing=True)

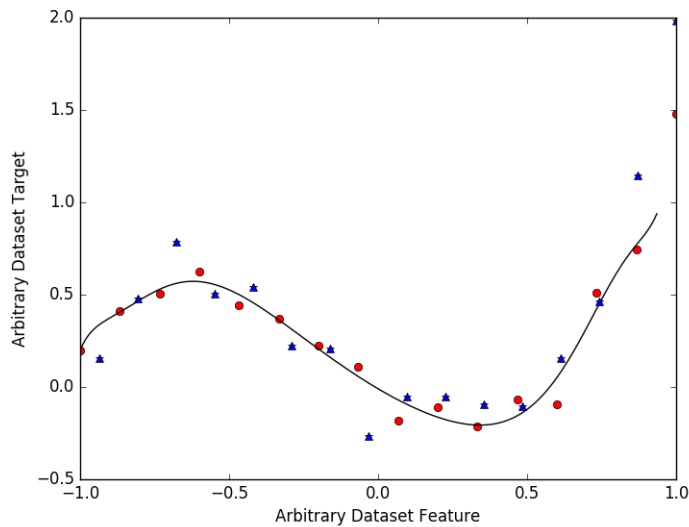
avg_rmse = np.sqrt(np.sum((np.dot(X,w) - y)**2)/len(y))
avg_rmse_test = np.sqrt(np.sum((np.dot(X_test, w) - y_test)**2)/len(
    y_test))

plt.plot(x,y,'ro')
plt.plot(xhat,yhat,'k-')
plt.plot(x_test,test_set[1],'b*')
plt.plot(x_test,y_test,'b^')
plt.xlabel('Arbitrary Dataset Feature')
plt.ylabel('Arbitrary Dataset Target')
plt.show()

if __name__ == "__main__":
    dataset_file = "P1C_training.csv"
    testset_file = "P1C_test.csv"
    dataset = pd.read_csv(dataset_file, header=None)
    testset = pd.read_csv(testset_file, header=None)
    fit_with_l2(15, 1e-4, dataset, testset)

```

1C



1D

```

import pandas as pd
import numpy as np

```

```

import matplotlib as mpl
import matplotlib.pyplot as plt

def fit_with_l2(degree, gamma, dataset, test_set):
    #data and target from pandas frame
    x = dataset[0].as_matrix().astype(float)
    y = dataset[1].as_matrix().astype(float)

    x_test = test_set[0].as_matrix().astype(float)
    y_test = test_set[1].as_matrix().astype(float)

    xhat = np.linspace(x.min(), x.max(), 200)
    Xhat = np.vander(xhat, N=degree + 1, increasing=True)
    x = 2*(dataset[0] - x.min()/(x.max()-x.min()) - 1) #this is x_tilde
    x_test = 2*(test_set[0] - x.min()/(x.max()-x.min()) - 1)

    # Vandermonde matrix (design matrix)
    X = np.vander(x, degree+1, increasing=True)

    # Identity matrix for regularization
    Eye = np.eye(X.shape[1])
    #Eye hat, to not regularize bias
    Eye[0, 0] = 0

    # Solve for weight's set (params) (training)
    w = np.linalg.solve(np.dot(X.T, X) + gamma*Eye, np.dot(X.T, y))

    yhat = np.dot(Xhat, w) # 250x2, 16x1

    X_test = np.vander(x_test, N=degree+1, increasing=True)

    avg_rmse = np.sqrt(np.sum((np.dot(X, w) - y)**2)/len(y))
    avg_rmse_test = np.sqrt(np.sum((np.dot(X_test, w) - y_test)**2)/len(
        y_test))
    return avg_rmse, avg_rmse_test

if __name__ == "__main__":
    dataset_file = "P1C_training.csv"
    testset_file = "P1C_test.csv"
    dataset = pd.read_csv(dataset_file, header=None)
    testset = pd.read_csv(testset_file, header=None)


    degs = []
    rmse = []
    rmse_test = []

    for i in range(1, 16):
        avg_rmse, avg_rmse_test = fit_with_l2(i, 0, dataset, testset)
        degs.append(i)
        rmse.append(avg_rmse)
        rmse_test.append(avg_rmse_test)

    plt.plot(degs, rmse, 'b-')
    plt.plot(degs, rmse_test, 'r-')
    plt.xlabel('Degree N')
    plt.ylabel('Average RMSE')

```

```
plt.show()
```



1E

```
import pandas as pd
import numpy as np
import matplotlib as mpl
import matplotlib.pyplot as plt

def fit_with_l2(degree, gamma, dataset, test_set):
    #data and target from pandas frame
    x = dataset[0].as_matrix().astype(float)
    y = dataset[1].as_matrix().astype(float)

    x_test = test_set[0].as_matrix().astype(float)
    y_test = test_set[1].as_matrix().astype(float)

    xhat = np.linspace(x.min(), x.max(), 200)
    Xhat = np.vander(xhat, N=degree + 1, increasing=True)
    x = 2*(dataset[0] - x.min())/(x.max()-x.min()) - 1 #this is x_tilde
    x_test = 2*(test_set[0] - x.min())/(x.max()-x.min()) - 1

    # Vandermonde matrix (design matrix)
    X = np.vander(x, degree+1, increasing=True)

    # Identity matrix for regularization
    Eye = np.eye(X.shape[1])
    #Eye hat, to not regularize bias
    Eye[0, 0] = 0

    # Solve for weight's set (params) (training)
```

```

w = np.linalg.solve(np.dot(X.T,X) + gamma*Eye,np.dot(X.T,y))

yhat = np.dot(Xhat,w) # 250x2, 16x1

X_test = np.vander(x_test, N=degree+1, increasing=True)

avg_rmse = np.sqrt(np.sum((np.dot(X,w) - y)**2)/len(y))
avg_rmse_test = np.sqrt(np.sum((np.dot(X_test, w) - y_test)**2)/len(
    y_test))
return avg_rmse, avg_rmse_test

if __name__ == "__main__":
    dataset_file = "P1C_training.csv"
    testset_file = "P1C_test.csv"
    dataset = pd.read_csv(dataset_file, header=None)
    testset = pd.read_csv(testset_file, header=None)


    rmse = []
    rmse_test = []

    gammas = np.logspace(-10, 3, 150)

    for i in gammas:
        avg_rmse, avg_rmse_test = fit_with_l2(15, i, dataset,
            testset)
        rmse.append(avg_rmse)
        rmse_test.append(avg_rmse_test)

    plt.semilogx(gammas,rmse,'b-')
    plt.semilogx(gammas, rmse_test, 'r-')
    plt.xlabel('Gamma')
    plt.ylabel('Average RMSE')
    plt.show()

```



1F

```

import pandas as pd
import numpy as np
import matplotlib as mpl
import matplotlib.pyplot as plt
from sklearn import linear_model

def fit_with_l1(degree, gamma, dataset, test_set):
    #data and target from padas frame
    x = dataset[0].as_matrix().astype(float)
    y = dataset[1].as_matrix().astype(float)

    x_test = test_set[0].as_matrix().astype(float)
    y_test = test_set[1].as_matrix().astype(float)

    xhat = np.linspace(x.min(), x.max(), 200)
    Xhat = np.vander(xhat, N=degree + 1, increasing=True)
    x = 2*(dataset[0] - x.min()/(x.max()-x.min())) - 1 #this is x_tilde
    x_test = 2*(test_set[0] - x.min()/(x.max()-x.min())) - 1

    # Vandermond matrix (design matrix)
    X = np.vander(x, degree+1, increasing=True)

    # Identity matrix for regularization
    Eye = np.eye(X.shape[1])
    #Eye hat, to not regularize bias
    Eye[0, 0] = 0

    # Solve for weight's set (params) (training)
    w = np.linalg.solve(np.dot(X.T, X) + gamma*Eye, np.dot(X.T, y))

    yhat = np.dot(Xhat, w) # 250x2, 16x1

    X_test = np.vander(x_test, N=degree+1, increasing=True)

    lasso = linear_model.Lasso(alpha=gamma, max_iter=10000000)
    lasso.fit(X, y)

    avg_rmse = np.sqrt(np.sum((np.dot(X, w) - y)**2)/len(y))
    avg_rmse_test = np.sqrt(np.sum((np.dot(X_test, w) - y_test)**2)/len(
        y_test))
    return avg_rmse, avg_rmse_test

if __name__ == "__main__":
    dataset_file = "P1C_training.csv"
    testset_file = "P1C_test.csv"
    dataset = pd.read_csv(dataset_file, header=None)
    testset = pd.read_csv(testset_file, header=None)

    rmses = []
    rmses_test = []

    gammas = np.logspace(-10, 3, 40)

    for i in gammas:


```

```

print("running lasso {}".format(i))
avg_rmse, avg_rmse_test = fit_with_l1(15, i, dataset,
    testset)
rmsees.append(avg_rmse)
rmsees_test.append(avg_rmse_test)

plt.semilogx(gammas,rmsees,'b-')
plt.semilogx(gammas, rmsees_test, 'r-')
plt.xlabel('Gamma')
plt.ylabel('Average RMSE')
plt.show()

```



The plots look nearly identical to me with the exception that 1E is smoother because i used less points in 1F (Lasso takes a long time to run with small alphas)

1G*

```

import pandas as pd
import numpy as np
import matplotlib as mpl
import matplotlib.pyplot as plt

def fit_with_l2(degree, gamma, dataset, test_set):
    #data and target from padas frame
    x = dataset[0].as_matrix().astype(float)
    y = dataset[1].as_matrix().astype(float)

    x_test = test_set[0].as_matrix().astype(float)

```



```

y_test = test_set[1].as_matrix().astype(float)

xhat = np.linspace(x.min(),x.max(), 200)
Xhat = np.vander(xhat, N=degree + 1, increasing=True)
x = 2*(dataset[0] - x.min()/(x.max()-x.min()) - 1 #this is x_tilde
x_test = 2*(test_set[0] - test_set[0].min()/(test_set[0].max()-
    test_set[0].min()) - 1

# Vandermonde matrix (design matrix)
X = np.polynomial.legendre.legvander(x,degree)

# Identity matrix for regularization
Eye = np.eye(X.shape[1])
#Eye hat, to not regularize bias
Eye[0, 0] = 0

# Solve for weight's set (params) (training)
w = np.linalg.solve(np.dot(X.T,X) + gamma*Eye,np.dot(X.T,y))

yhat = np.dot(Xhat,w) # 250x2, 16x1

X_test = np.polynomial.legendre.legvander(x_test, degree)

avg_rmse = np.sqrt(np.sum((np.dot(X,w) - y)**2)/len(y))
avg_rmse_test = np.sqrt(np.sum((np.dot(X_test, w) - y_test)**2)/len(
    y_test))
return avg_rmse, avg_rmse_test


if __name__ == "__main__":
    dataset_file = "P1C_training.csv"
    testset_file = "P1C_test.csv"
    dataset = pd.read_csv(dataset_file, header=None)
    testset = pd.read_csv(testset_file, header=None)

    degs = []
    rmses = []
    rmses_test = []

    for i in range(1, 21):
        avg_rmse, avg_rmse_test = fit_with_l2(i, 0, dataset, testset
        )
        degs.append(i)
        rmses.append(avg_rmse)
        rmses_test.append(avg_rmse_test)

    plt.plot(degs,rmses,'b-')
    plt.plot(degs, rmses_test, 'r-')
    plt.xlabel('Degree N')
    plt.ylabel('Average RMSE')
    plt.show()

```



Probability and Bayes Theorem

2A Bayesian Inference

$$P(D) = 10^{-4}$$

$$P(\neg D) = 1 - 10^{-4}$$

$$P(T|D) = 0.99$$

$$P(T) = P(D)P(T|D) + P(\neg D)P(\neg T|D)$$

$$P(D|T) = \frac{P(T|D)P(D)}{P(T)}$$

$$P(D|T) = \frac{P(T|D)P(D)}{P(T)}$$

$$P(D|T) = \frac{P(T|D)P(D)}{P(D)P(T|D) + P(\neg D)P(\neg T|D)}$$

$$P(D|T) = \frac{0.99 \cdot 10^{-4}}{10^{-4} \cdot 0.99 + (1 - 10^{-4}) \cdot 0.01}$$

$$P(D|T) = 0.0098039$$

$$P(T_N|D) = 0.99999$$

$$P(D|T_N) = \frac{0.99999 \cdot 10^{-4}}{10^{-4} \cdot 0.99999 + (1 - 10^{-4}) \cdot 0.00001}$$

$$P(D|T) = 0.9090983$$

2B Bayesian Linear Regression

2C* Updating Bayesian Linear Regression

Naive Bayes

3A

We do not need to compute the denominator in the classification rule for Naive Bayes because we are assuming the data is independent and identically distributed. That is, it is a normal distribution with no dependence between features.

3B

```
import pandas as pd
import numpy as np
import matplotlib as mpl
import matplotlib.pyplot as plt

from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix

if __name__ == "__main__":
    digits = load_digits()
    data = np.round(digits['data']/16.0)

    X = data      # n x m matrix of features
    y = digits.target # n vector of classes
    X,X_test,y,y_test = train_test_split(X,y,test_size=0.33,random_state
        =42) # Split into 33% test and 67% training sets

    classes = np.linspace(0,9,10)

    m = X.shape[0] # Number of data instances
    m_test = X_test.shape[0] # Number of test data instances
    N = 10 # Number of classes
    n = X.shape[1] # Number of features

    mu_array = np.zeros((n,N))
    sigma2_array = np.zeros((n,N))
    prior_array = np.zeros((N))

    #Learning phase
    for k in range(N): #Loop over each class label
        C_k = classes[k]
```

```

prior = sum(y==C_k)/float(y.shape[0])
                                # Count the number of data
                                where the label is C_k
mu = np.sum(X[y==C_k],axis=0)/len(X[y==C_k])
                                # Take the mean of those features
                                where the corresponding label is C_k
mu_array[:,k] = mu
                                # Store in
                                the arrays we created above
prior_array[k] = prior

class_probabilities = np.zeros((m,N)) # The probabilities for
for i,x in enumerate(X): # Loop over the training data instances
    for k in range(N): # Loop over the classes
        prior = prior_array[k]
        mu = mu_array[:,k]
        # likelihood = np.prod(np.exp(-(x-mu)**2/(2*sigma2)))
        ) #change me
        likelihood = np.prod(np.power(mu, x) * np.power((1-
            mu), 1-x) )
        posterior_k = prior*likelihood
        class_probabilities[i,k] = posterior_k

class_probabilities /= np.sum(class_probabilities,axis=1,keepdims=
    True)

y_pred_train = np.argmax(class_probabilities,axis=1)

cm_train = confusion_matrix(y,y_pred_train)
print(cm_train)
print("training accuracy:", 1 - sum(abs(y!=y_pred_train))/float(m))

# Test set predictions
class_probabilities = np.zeros((m_test,N))

for i,x in enumerate(X_test):
    for k in range(N):
        prior = prior_array[k]
        mu = mu_array[:,k]
        sigma2 = sigma2_array[:,k]
        likelihood = np.prod(np.power(mu, x) * np.power((1-
            mu), 1-x) )
        posterior_k = prior*likelihood
        class_probabilities[i,k] = posterior_k

class_probabilities /= class_probabilities.sum(axis=1,keepdims=True)
y_pred_test = np.argmax(class_probabilities,axis=1)

cm_test = confusion_matrix(y_test,y_pred_test)
print(cm_test)
print("test accuracy:", 1 - sum(abs(y_test-y_pred_test))/float(
    m_test))

```

training confusion matrix:

$$\begin{bmatrix} 119 & 0 & 0 & 0 & 1 & 2 & 0 & 0 & 0 & 1 \\ 0 & 108 & 4 & 0 & 0 & 1 & 1 & 0 & 7 & 6 \\ 1 & 2 & 116 & 1 & 0 & 0 & 0 & 0 & 2 & 3 \\ 0 & 2 & 1 & 111 & 0 & 2 & 0 & 6 & 2 & 3 \\ 0 & 1 & 0 & 0 & 108 & 0 & 1 & 4 & 3 & 0 \\ 0 & 0 & 1 & 2 & 1 & 99 & 0 & 0 & 0 & 6 \\ 0 & 3 & 0 & 0 & 1 & 0 & 119 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 115 & 1 & 0 \\ 0 & 11 & 3 & 0 & 0 & 3 & 0 & 2 & 101 & 2 \\ 0 & 0 & 0 & 2 & 0 & 2 & 0 & 3 & 2 & 103 \end{bmatrix}$$

training accuracy: 0.913549459684

testing confusion matrix:

$$\begin{bmatrix} 52 & 0 & 0 & 0 & 2 & 1 & 0 & 0 & 0 & 0 \\ 0 & 37 & 5 & 0 & 1 & 0 & 0 & 0 & 7 & 5 \\ 0 & 4 & 46 & 0 & 0 & 0 & 0 & 0 & 2 & 0 \\ 0 & 0 & 2 & 45 & 0 & 0 & 0 & 1 & 3 & 5 \\ 0 & 0 & 0 & 0 & 63 & 0 & 0 & 1 & 0 & 0 \\ 2 & 0 & 0 & 0 & 1 & 60 & 0 & 0 & 0 & 10 \\ 0 & 0 & 0 & 0 & 0 & 1 & 55 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 & 1 & 0 & 58 & 0 & 1 \\ 0 & 2 & 2 & 1 & 0 & 1 & 0 & 0 & 45 & 1 \\ 0 & 0 & 0 & 4 & 0 & 1 & 0 & 4 & 2 & 57 \end{bmatrix}$$

test accuracy: 0.872053872054