

CSCI547 Machine Learning

Homework 2

Zachary Falkner
Department of Computer Science
University of Montana

March 9, 2018

1 Logistic Regression

1A

```
import numpy as np
import pandas as pd

# The sigmoid function
def _sigmoid(w,X):
    z = np.dot(X,w)
    return 1./(1+np.exp(-z))

# The objective function
def _J_fun(w,X,y):
    return -sum(y*np.log(_sigmoid(w,X)) + (1-y)*np.log(1-_sigmoid(w,X)))

# The gradient of the objective function
def _gradient_fun(w,X,y):
    return np.dot(_sigmoid(w,X)-y,X)

class LogisticRegression:
    def __init__(self, eta=None, epochs=10000, w=None):
        self.eta= eta
        self.epochs = epochs
        self.w = w
        self.cost_over_epochs = []
```

```

def fit(self, x, y):
    N = len(y)

    for i in range(self.epochs):
        grad_w = _gradient_fun(self.w,x,y)      # Compute the gradient
                                                of the objective function
        self.w -= np.dot(self.eta,grad_w)
        self.cost_over_epochs.append(_J_fun(self.w,x,y))

        classification_error = sum((_sigmoid(self.w,x)>0.5)==y)/N
    return classification_error

def score(self, x, y):
    N = len(y)

    classification_error = sum((_sigmoid(self.w,x)>0.5)==y)/N
    return classification_error

```

```

import numpy as np
import pandas as pd

from logistic_regression import LogisticRegression

if __name__ == '__main__':
    DATAFILE = 'lobster_survive.dat'
    df = pd.read_csv(DATAFILE,header=0, sep=r"\s{2,}")
    x = df.iloc[:, 0].as_matrix().astype(float)
    y = df.iloc[:, 1].as_matrix().astype(float)
    x = np.vander(x,2,increasing=True)

    #learning rate, tensor
    eta = np.array([[0.000001,0],[0,0.000000001]])

    #number of iterations
    epochs = 200000

    #weights
    w = np.array([-1.,0.5])

    lr = LogisticRegression(eta=eta, epochs=epochs, w=w)
    error = lr.fit(x, y)
    print("classification error: {}".format(error))

```

classification error: 0.710691823899371

1B

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from scipy.stats import zscore

```

```

from logistic_regression import LogisticRegression

if __name__ == '__main__':
    TRAINFILE = 'titanic_train.csv'
    TESTFILE = 'titanic_test.csv'
    train_df = pd.read_csv(TRAINFILE, header=0)
    test_df = pd.read_csv(TESTFILE, header=0)

    x_data = train_df.iloc[:, 2:]

    # drop useless columns
    x_data.drop(['Cabin', 'Ticket', 'Name'], axis=1, inplace=True)

    # one hot sex and embarked
    x_data = pd.get_dummies(x_data, columns=['Sex', 'Embarked'])

    # fill missing age data with mean...
    x_data['Age'] = x_data['Age'].fillna(x_data['Age'].mean())

    # normalize with zscores
    numeric_cols = ['Parch', 'SibSp', 'Age', 'Fare']
    x_data[numeric_cols] = x_data[numeric_cols].apply(zscore)

    x = x_data.as_matrix().astype(float)
    y = train_df.iloc[:, 1].as_matrix().astype(float)

    x_data_test = test_df.iloc[:, 2:]

    # drop useless columns
    x_data_test.drop(['Cabin', 'Ticket', 'Name'], axis=1, inplace=True)

    # one hot sex and embarked
    x_data_test = pd.get_dummies(x_data_test, columns=['Sex', 'Embarked',
    ])

    # fill missing age data with mean...
    x_data_test['Age'] = x_data_test['Age'].fillna(x_data_test['Age'].
    mean())

    # normalize with zscores
    x_data_test[numeric_cols] = x_data_test[numeric_cols].apply(zscore)

    x_test = x_data_test.as_matrix().astype(float)
    y_test = test_df.iloc[:, 1].as_matrix().astype(float)

    # learning rate, tensor
    eta = np.eye(10) * np.array([0.000001] * 10)

    # number of iterations
    epochs = 200000

    # weights
    w = np.array([-1., 0.5, 0.5, 0.5, 0.5, 0.5, 0.5, 0.5, 0.5, 0.5])

```

```

lr = LogisticRegression(eta=eta, epochs=epochs, w=w)

error = lr.fit(x, y)
print("training error: {}".format(error))

test_error = lr.score(x_test, y_test)
print("test error: {}".format(test_error))

iterations = len(lr.cost_over_epochs)
print("training iterations: {}".format(iterations))


plt.plot(lr.cost_over_epochs, 'k')
plt.xlabel('training iterations N')
plt.ylabel('cost')
plt.show()

```

training error: 0.7962962962962963

test error: 0.8047138047138047

training iterations: 100000



1C*

```

import numpy as np
import pandas as pd

# The sigmoid function
def _sigmoid(w,X):
    z = np.dot(X,w)
    return 1./(1+np.exp(-z))

```

```

# The objective function
def _J_fun(w,X,y):
    return -sum(y*np.log(_sigmoid(w,X)) + (1-y)*np.log(1-_sigmoid(w,X)))

# The gradient of the objective function
def _gradient_fun(w,X,y):
    return np.dot(_sigmoid(w,X)-y,X)

class LogisticRegression:
    def __init__(self, eta=None, epochs=10000, w=None, enable_early_stop
        =False, early_stop_tolerance=10):
        self.eta= eta
        self.epochs = epochs
        self.w = w
        self.cost_over_epochs = []
        self.gradient_over_epochs = []

        self.enable_early_stop = enable_early_stop
        self.early_stop_tolerance = early_stop_tolerance

    def fit(self, x, y):
        N = len(y)

        for i in range(self.epochs):
            grad_w = _gradient_fun(self.w,x,y) # Compute the
            gradient of the objective function
            self.w -= np.dot(self.eta,grad_w)
            self.cost_over_epochs.append(_J_fun(self.w,x,y))
            self.gradient_over_epochs.append(grad_w)

            if(self.shouldTerminateEarly()):
                break;

        classification_error = sum((_sigmoid(self.w,x)>0.5)==y)/N
        return classification_error

    def score(self, x, y):
        N = len(y)

        classification_error = sum((_sigmoid(self.w,x)>0.5)==y)/N
        return classification_error

    def shouldTerminateEarly(self):
        if(self.enable_early_stop == True):
            grads = self.gradient_over_epochs[-1]
            return (all(g < self.early_stop_tolerance for g in
                grads)
                and all(g > -1 * self.early_stop_tolerance for g in
                grads) )
            return False

```

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from scipy.stats import zscore

from logistic_regression import LogisticRegression

if __name__ == '__main__':
    TRAINFILE = 'titanic_train.csv'
    TESTFILE = 'titanic_test.csv'
    train_df = pd.read_csv(TRAINFILE, header=0)
    test_df = pd.read_csv(TESTFILE, header=0)

    x_data = train_df.iloc[:, 2:]

    # drop useless columns
    x_data.drop(['Cabin', 'Ticket', 'Name'], axis=1, inplace=True)

    # one hot sex and embarked
    x_data = pd.get_dummies(x_data, columns=['Sex', 'Embarked'])

    # fill missing age data with mean...
    x_data['Age'] = x_data['Age'].fillna(x_data['Age'].mean())

    # normalize with zscores
    numeric_cols = ['Parch', 'SibSp', 'Age', 'Fare']
    x_data[numeric_cols] = x_data[numeric_cols].apply(zscore)

    x = x_data.as_matrix().astype(float)
    y = train_df.iloc[:, 1].as_matrix().astype(float)

    x_data_test = test_df.iloc[:, 2:]

    # drop useless columns
    x_data_test.drop(['Cabin', 'Ticket', 'Name'], axis=1, inplace=True)

    # one hot sex and embarked
    x_data_test = pd.get_dummies(x_data_test, columns=['Sex', 'Embarked'])

    # fill missing age data with mean...
    x_data_test['Age'] = x_data_test['Age'].fillna(x_data_test['Age'].mean())

    # normalize with zscores
    x_data_test[numeric_cols] = x_data_test[numeric_cols].apply(zscore)

    x_test = x_data_test.as_matrix().astype(float)
    y_test = test_df.iloc[:, 1].as_matrix().astype(float)

    # learning rate, tensor
    eta = np.eye(10) * np.array([0.000001] * 10)

    # number of iterations
    epochs = 100000

```

```

#weights
w = np.array([-1.,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5])

lr = LogisticRegression(eta=eta, epochs=epochs, w=w, enable_early_stop=True,
    early_stop_tolerance=10)

error = lr.fit(x, y)
print("training error: {}".format(error))

test_error = lr.score(x_test, y_test)
print("test error: {}".format(test_error))

iterations = len(lr.cost_over_epochs)
print("training iterations: {}".format(iterations))

plt.plot(lr.cost_over_epochs, 'k')
plt.plot(lr.gradient_over_epochs, 'r')
plt.xlabel('training iterations N')
plt.ylabel('gradients, cost')
plt.show()


```

training error: 0.7962962962962963

test error: 0.8080808080808081

training iterations: 45137

for this i plotted cost and gradients over iterations and noticed that when the gradients all converge towards 0 the cost is converging. seems sound, got similar results with less than 1/4 the iterations.



2 Neural Networks

2A

Ok, Done!

2B

```
from __future__ import division, print_function

import numpy as np

class Network(object):
    """
    Neural network for softmax regression problems
    """

    def __init__(self, layer_number_of_nodes, layer_activation_functions,
                  layer_has_bias, layer_weight_means_and_stds=None):
        self.layer_number_of_nodes = layer_number_of_nodes          # Of
                               nodes in each layer
        self.layer_activation_functions = [None]
        # Add an identity activation function here!
        for act in layer_activation_functions:
            if act=='softmax':
                self.layer_activation_functions.append(self._softmax
                )
            if act=='sigmoid':
                self.layer_activation_functions.append(self._sigmoid
                )
            if act=='leaky_relu':
                self.layer_activation_functions.append(self.
                _leaky_relu)
            if act=='gaussian':
                self.layer_activation_functions.append(self.
                _gaussian)
            if act=='identity':
                self.layer_activation_functions.append(self.
                _identity)

        self.layer_has_bias = layer_has_bias                        #
                               Whether to add a bias node to each layer
        self.L = len(self.layer_number_of_nodes)                    #
                               Number of layers

        self.weights = [np.array([])]

        # Create arrays to hold the weights, which are N_l(+1) by N_(l+1)
        for i in range(self.L-1):
            # if we have a normal distribution and standard deviation,
            # then generate random weights from that distribution,
            if layer_weight_means_and_stds is not None:
                w = layer_weight_means_and_stds[i][1]*np.random.
                    randn(self.layer_number_of_nodes[i] + self.
```



```

        layer_has_bias[i],self.layer_number_of_nodes[i
        +1]) + layer_weight_means_and_stds[i][0]
        # Otherwise just initialize the weights to zero
    else:
        w = np.zeros((self.layer_number_of_nodes[i] + self.
            layer_has_bias[i],self.layer_number_of_nodes[i
            +1]))
        self.weights.append(w)

def feed_forward(self, feature):
    # evaluate the neural network for a vector-valued input
    m = feature.shape[0]

    # Append a column of ones to the input if a bias is desired
    if self.layer_has_bias[0]:
        z = np.column_stack((np.ones((m)), feature))
    else:
        z = feature

    # Initialize lists to hold the node inputs and outputs, treating the
    # input values as the output of the first node
    self.a_vals = [None]
    self.z_vals = [z]

    # Loop over the remaining layers
    for l in range(1, self.L):
        # Take the linear combination of the previous layers outputs
        # ( $z^{(l-1)}$ ) and weights ( $w^{(l)}$ ) to form  $a^{(l)}$ 
        a = np.dot(self.z_vals[l-1], self.weights[l])
        # Run a through the activation function to form  $z^{(l)}$ 
        z = self.layer_activation_functions[l](a)
        # If a bias is desired, append a column of ones to z
        if self.layer_has_bias[l]:
            z = np.column_stack((np.ones((m)), z))
        # Store these values (for computing the gradient later)
        self.a_vals.append(a)
        self.z_vals.append(z)

    return z

def _J_fun(self, feature, label):
    # Add your sum square error evaluation here!
    if self.layer_activation_functions[-1]==self._identity:
        cost_function_data = np.sum(np.sum((1/2) * np.power(label -
            self.feed_forward(feature), 2), axis=1), axis=0)
    if self.layer_activation_functions[-1]==self._softmax:
        # Model objective function -- Cross-entropy
        cost_function_data = -np.sum(np.sum(label*np.log(self.
            feed_forward(feature)), axis=1), axis=0)
    else:
        print('Only softmax supported for final layer')

    # Add regularization here!
    # TODO
    cost_function_reg = 0
    return cost_function_data + cost_function_reg

def _gradient_fun(self, feature, label):

```

```

# Compute the gradient via backpropagation
m = feature.shape[0]

# Initialize gradient arrays (same shape as the weights)
grads = [np.zeros_like(w) for w in self.weights]

# Compute dJ/da (aka the delta term) for the final layer. This
# often involves
# Some algebraic simplification when cost function is selected
# judiciously, so
# this is coded by hand here.

l = self.L-1 #last layer

z = self.z_vals[l] # Current layer out
z_previous = self.z_vals[l-1] # Last layer out
a = self.a_vals[l] # Current layer in
w = self.weights[l] # Last layer weights
activation = self.layer_activation_functions[l] # Current layer
activation
if activation==self._softmax or activation==self._identity:
    delta_l = (z - label) # Current layer
    error
# Add gradient of SSE here!
else:
    print('Only softmax and identity supported for final layer')

grads[l] = np.dot(z_previous.T,delta_l) # gradient due to data
misfit

# Add gradient of regularization here!
model_norm_gradient = 0
grads[l] += model_norm_gradient # add gradient
due to regularization

# Loop over the remaining layers
for l in range(self.L-2,0,-1):

    z_previous = self.z_vals[l-1] # last
    layer output
    a = self.a_vals[l] # Current
    layer input

    w_next = self.weights[l+1][1:] # weights from the next layer
    , excluding bias weights
    activation = self.layer_activation_functions[l] # Current
    layer activation

    delta_l = np.dot(delta_l,w_next.T)*activation(a,dx=1) #
    Current layer error
    grads[l] = np.dot(z_previous.T,delta_l) # Gradient due to
    data misfit

    # Add gradient of regularization here!
    model_norm_gradient = 0
    grads[l] += model_norm_gradient # add gradient
    due to regularization

```

```

        return grads

    @staticmethod
    def _softmax(X,dx=0):
        if dx==0:
            return np.exp(X)/np.repeat(np.sum(np.exp(X),axis=1,keepdims=True),X.shape[1],axis=1)
    @staticmethod
    def _sigmoid(X,dx=0):
        if dx==0:
            return 1./(1+np.exp(-X))
        if dx==1:
            s = 1./(1+np.exp(-X))
            return s*(1-s)

    @staticmethod
    def _leaky_relu(X,dx=0):
        if dx==0:
            return (X>0)*X + 0.01*(X<=0)*X
        if dx==1:
            return (X>0) + 0.01*(X<=0)

    @staticmethod
    def _gaussian(X,dx=0):
        if dx==0:
            return np.exp(-X**2)
        if dx==1:
            return -2*X*np.exp(-X**2)

    # Add @staticmethod for identity here
    @staticmethod
    def _identity(X,dx=0):
        if dx == 0:
            return X
        return np.onelike(X)

```

```

import numpy as np
import matplotlib.pyplot as plt

from neural_network import Network

if __name__ == '__main__':
    n = 1
    m = 100
    N = 1
    X = np.random.rand(m, 1)
    y = np.exp(-np.sin(np.power(X, 3) * 4 * np.pi))

    X_test = np.random.rand(m, 1)
    y_test = np.exp(-np.sin(np.power(X_test, 3) * 4 * np.pi))

    nn = Network([n,20,N],[None,'sigmoid','identity'],[True,True,False],
        layer_weight_means_and_stds=[[0,0.1),(0,0.1)])

    eta = 0.001

```

```

N_iterations = 10000

T = y

# Perform gradient descent
for i in range(N_iterations):

    # For stochastic gradient descent, take random samples of X
    # and T

    # Run the features through the neural net (to compute a and
    # z)
    y_pred = nn.feed_forward(X)

    # Compute the gradient
    grad_w = nn._gradient_fun(X,T)

    # Update the neural network weight matrices
    for w,gw in zip(nn.weights,grad_w):
        w -= eta*gw

    # Print some statistics every thousandth iteration
    if i%1000==0:
        misclassified = sum(np.argmax(y_pred,axis=1)!=y.
            ravel())
        print ("Iteration: {0}, Objective Function Value:
            {1:3f}, Misclassified: {2}".format(i,nn._J_fun(X
            ,T), misclassified))

    # Predict the training data and classify
    y_pred = np.argmax(nn.feed_forward(X_test),axis=1)
    print ("Test data accuracy: {0:3f}".format(1-sum(y_pred!=y_test.
        ravel())/float(len(y_test))))

```

2C

```

from __future__ import division,print_function

import numpy as np

class Network(object):
    """
    Neural network for softmax regression problems
    """

    def __init__(self,layer_number_of_nodes,layer_activation_functions,
        layer_has_bias,layer_weight_means_and_std=None, gama=0.1,
        regularization=None):
        self.layer_number_of_nodes = layer_number_of_nodes # Of
            nodes in each layer
        self.layer_activation_functions = [None]
        # Add an identity activation function here!
        for act in layer_activation_functions:

```

```

        if act=='softmax':
            self.layer_activation_functions.append(self._softmax
            )
        if act=='sigmoid':
            self.layer_activation_functions.append(self._sigmoid
            )
        if act=='leaky_relu':
            self.layer_activation_functions.append(self.
            _leaky_relu)
        if act=='gaussian':
            self.layer_activation_functions.append(self.
            _gaussian)
        if act=='identity':
            self.layer_activation_functions.append(self.
            _identity)

self.layer_has_bias = layer_has_bias #
    Whether to add a bias node to each layer
self.L = len(self.layer_number_of_nodes) #
    Number of layers
self.gama = gama

self.weights = [np.array([])]
self.regularization = regularization

# Create arrays to hold the weights, which are N_l(+1) by N_(l+1)
for i in range(self.L-1):
    # if we have a normal distribution and standard deviation,
    then generate random weights from that distribution,
    if layer_weight_means_and_stds is not None:
        w = layer_weight_means_and_stds[i][1]*np.random.
            randn(self.layer_number_of_nodes[i] + self.
            layer_has_bias[i],self.layer_number_of_nodes[i
            +1]) + layer_weight_means_and_stds[i][0]
    # Otherwise just initialize the weights to zero
    else:
        w = np.zeros((self.layer_number_of_nodes[i] + self.
            layer_has_bias[i],self.layer_number_of_nodes[i
            +1]))
    self.weights.append(w)

def feed_forward(self,feature):
    # evaluate the neural network for a vector-valued input
    m = feature.shape[0]

    # Append a column of ones to the input if a bias is desired
    if self.layer_has_bias[0]:
        z = np.column_stack((np.ones((m)),feature))
    else:
        z = feature

    # Initialize lists to hold the node inputs and outputs, treating the
    input values as the output of the first node
    self.a_vals = [None]
    self.z_vals = [z]

# Loop over the remaining layers

```

```

        for l in range(1,self.L):
            # Take the linear combination of the previous layers outputs
            # (z^(l-1)) and weights (w^(l)) to form a^(l)
            a = np.dot(self.z_vals[l-1],self.weights[l])
            # Run a through the activation function to form z^(l)
            z = self.layer_activation_functions[l](a)
            # If a bias is desired, append a column of ones to z
            if self.layer_has_bias[l]:
                z = np.column_stack((np.ones((m)),z))
            # Store these values (for computing the gradient later)
            self.a_vals.append(a)
            self.z_vals.append(z)
        return z

def _J_fun(self,feature,label):
    # Add your sum square error evaluation here!
    if self.layer_activation_functions[-1]==self._identity:
        cost_function_data = np.sum(np.sum((1/2) * np.power(label -
            self.feed_forward(feature), 2), axis=1), axis=0)
    if self.layer_activation_functions[-1]==self._softmax:
        # Model objective function -- Cross-entropy
        cost_function_data = -np.sum(np.sum(label*np.log(self.
            feed_forward(feature)),axis=1),axis=0)
    else:
        print('Only softmax supported for final layer')

    # Add regularization here!
    # TODO
    cost_function_reg = 0
    if(self.regularization == 'L1'):
        for w in self.weights:
            cost_function_reg = self.gama * np.sum(np.abs(w))
    if(self.regularization == 'L2'):
        for w in self.weights:
            cost_function_reg = self.gama * np.sum(np.power(w,
                2))

    return cost_function_data + cost_function_reg

def _gradient_fun(self,feature,label):
    # Compute the gradient via backpropagation
    m = feature.shape[0]

    # Initialize gradient arrays (same shape as the weights)
    grads = [np.zeros_like(w) for w in self.weights]

    # Compute dJ/da (aka the delta term) for the final layer. This
    # often involves
    # Some algebraic simplification when cost function is selected
    # judiciously, so
    # this is coded by hand here.

    l = self.L-1 #last layer

    z = self.z_vals[l] # Current layer out
    z_previous = self.z_vals[l-1] # Last layer out
    a = self.a_vals[l] # Current layer in

```

```

w = self.weights[l] # Last layer weights
activation = self.layer_activation_functions[l] # Current layer
activation
if activation==self._softmax or activation==self._identity:
    delta_l = (z - label) # Current layer
    error
# Add gradient of SSE here!
else:
    print('Only softmax and identity supported for final layer')

grads[l] = np.dot(z_previous.T,delta_l) # gradient due to data
misfit

# Add gradient of regularization here!
model_norm_gradient = 0
if(self.regularization == 'L1'):
    model_norm_gradient = self.gama * np.sign(w)
if(self.regularization == 'L2'):
    model_norm_gradient = self.gama * w
grads[l] += model_norm_gradient # add gradient
due to regularization

# Loop over the remaining layers
for l in range(self.L-2,0,-1):

    z_previous = self.z_vals[l-1] # last
    layer output
    a = self.a_vals[l] # Current
    layer input

    w_next = self.weights[l+1][1:] # weights from the next layer
    , excluding bias weights
    activation = self.layer_activation_functions[l] # Current
    layer activation

    delta_l = np.dot(delta_l,w_next.T)*activation(a,dx=1) #
    Current layer error
    grads[l] = np.dot(z_previous.T,delta_l) # Gradient due to
    data misfit

    # Add gradient of regularization here!
    model_norm_gradient = 0
    if(self.regularization == 'L1'):
        model_norm_gradient = self.gama * np.sign(self.
        weights[l])
    if(self.regularization == 'L2'):
        model_norm_gradient = self.gama * self.weights[l]
    grads[l] += model_norm_gradient # add gradient
    due to regularization

return grads

@staticmethod
def _softmax(X,dx=0):
    if dx==0:
        return np.exp(X)/np.repeat(np.sum(np.exp(X),axis=1,keepdims=
        True),X.shape[1],axis=1)

```

```

@staticmethod
def _sigmoid(X,dx=0):
    if dx==0:
        return 1./(1+np.exp(-X))
    if dx==1:
        s = 1./(1+np.exp(-X))
        return s*(1-s)

@staticmethod
def _leaky_relu(X,dx=0):
    if dx==0:
        return (X>0)*X + 0.01*(X<=0)*X
    if dx==1:
        return (X>0) + 0.01*(X<=0)

@staticmethod
def _gaussian(X,dx=0):
    if dx==0:
        return np.exp(-X**2)
    if dx==1:
        return -2*X*np.exp(-X**2)

# Add @staticmethod for identity here
@staticmethod
def _identity(X,dx=0):
    if dx == 0:
        return X
    return np.onelike(X)

```

2D

I overextended my schedule this week and did not get to complete 2D or 2E.

2E*

...

3 TensorFlow

3A

```

import matplotlib.pyplot as plt

...

# You can acquire the values of your layer weights with
w = sess.run(W_0)
w = w.reshape(10, 28, 28)


```



```
fig, axes = plt.subplots(2, 5, subplot_kw={'xticks': [], 'yticks': []})

for ax, x in zip(axes.flat, w):
    ax.imshow(x, interpolation=None, cmap='viridis')

plt.show()
```



3B

```
import argparse
import sys
import matplotlib.pyplot as plt

from tensorflow.examples.tutorials.mnist import input_data

import tensorflow as tf

# Tensorflow has the mnist data builtin
data_dir = '/tmp/tensorflow/mnist/input_data'

# Import data
mnist = input_data.read_data_sets(data_dir, one_hot=True)

n = 784 # Number of input features
N = 10  # Number of classes

# X is the vector of inputs (though it's just a placeholder until a
# tensorflow session is started)
X = tf.placeholder("float", [None, n])
```

```

# y is the vector of targets
y = tf.placeholder("float", [None, N])

# Create the model
# layer 1 weights and biases
W_0 = tf.Variable(tf.random_normal([n,N], stddev=0.01))
b_0 = tf.Variable(tf.random_normal([N], stddev=0.01))

# Create neural network
def multilayer_perceptron(x):
    out_layer = tf.add(tf.matmul(x,W_0),b_0)
    hidden_layer = tf.layers.dense(inputs=out_layer, units=300,
        activation=tf.nn.sigmoid)
    logits = tf.layers.dense(inputs=hidden_layer, units=10)
    return logits

# define prediction object
y_pred = multilayer_perceptron(X)

# Define loss function (combined softmax and cross-entropy output)
loss_op = tf.reduce_sum(tf.nn.softmax_cross_entropy_with_logits(logits=
    y_pred, labels=y))

# Specify learning rate
learning_rate = 0.001

# Define optimization step
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate)

# The optimization procedure (minimizing the softmax cross_entropy)
train_op = optimizer.minimize(loss_op)

# Initialize all the variables
# (tensorflow doesn't compute variable values unless run by a session)
sess = tf.InteractiveSession()
init = tf.global_variables_initializer().run()

# Train
N_iterations = 200000
sample_size = 10

for i in range(N_iterations):

    # Pull a sample from the training set
    batch_xs, batch_ys = mnist.train.next_batch(sample_size)

    # Run tensor flow objects: train_op updates the weights, loss_op
    # compute the cost function
    _,c = sess.run([train_op,loss_op], feed_dict={X: batch_xs, y:
        batch_ys})

    # Print statistics every 1000 steps
    if i%1000==0:
        # Test trained model
        pred = tf.nn.softmax(y_pred)

```

```

        correct_prediction = tf.equal(tf.argmax(pred, 1), tf.argmax(
            y,1))
        accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.
            float32))
        print("Accuracy:", accuracy.eval({X: mnist.test.images, y:
            mnist.test.labels}),c)

# You can acquire the values of your layer weights with
# w = sess.run(W_0)
# w = w.reshape(10, 28, 28)

# fig, axes = plt.subplots(2, 5, subplot_kw={'xticks': [], 'yticks': []})

# for ax, x in zip(axes.flat, w):
#     ax.imshow(x, interpolation=None, cmap='viridis')

# plt.show()

```